

# Project - Predictive Maintenance

September 15, 2018

## 1 A predictive maintenance project in the solar industry

- <li>1 - Context of problematic</li>
- <li>2 - Presentation of data</li>
- <li>3 - Study of hypothesis</li>
- <li>4 - Modeling and learning</li>
- <li>5 - Results and operation</li>

### 1.1 1 - Context of problematic

In the solar industry to make a photovoltaic panel, one of the crucial steps is the cutting of silicon ingots into wafers.

For this, diamond wire saws are used. These equipment are very efficient but also subject to drift manufacturing. Quality monitoring is therefore imperative. ##### In our context, important and regular drifts were observed. This is why a predictive maintenance prediction project was set up to reduce the impact of these production anomalies.

##### The aim was to predict when a drift would occur in order to realize an action of maintenance to prevent the anomaly (for example change the wire of the saw or recalibrate settings...).

### 1.2 2 - Presentation of data

For this project, we had a dataset with a dozen of variables. Each hour for each day (the factory was working on unceasing full time shift - 24 hours a day), when a cut is performed on a silicon ingot, a log file is produced with many useful informations.

Here are the most useful variables which be used for this project :

```
In [40]: import pandas as pd
import numpy as np

path = './Data'
fileName = 'Data_Maintenance.csv'

filePath = '///'.join([path,fileName])

data_maintenance = pd.read_csv(filePath,delimiter=';')
```

```
In [41]: data_maintenance.head(5)
```

```
Out[41]:
```

	Date	Time	DateTime	TTV	Thickness	SawGroove	SawEdge	\
0	01/01/2017	00:00	01/01/2017 00:00	30	250	120	1	
1	01/01/2017	01:00	01/01/2017 01:00	30	251	121	0	
2	01/01/2017	02:00	01/01/2017 02:00	30	251	120	0	
3	01/01/2017	03:00	01/01/2017 03:00	31	251	119	2	
4	01/01/2017	04:00	01/01/2017 04:00	32	251	124	1	

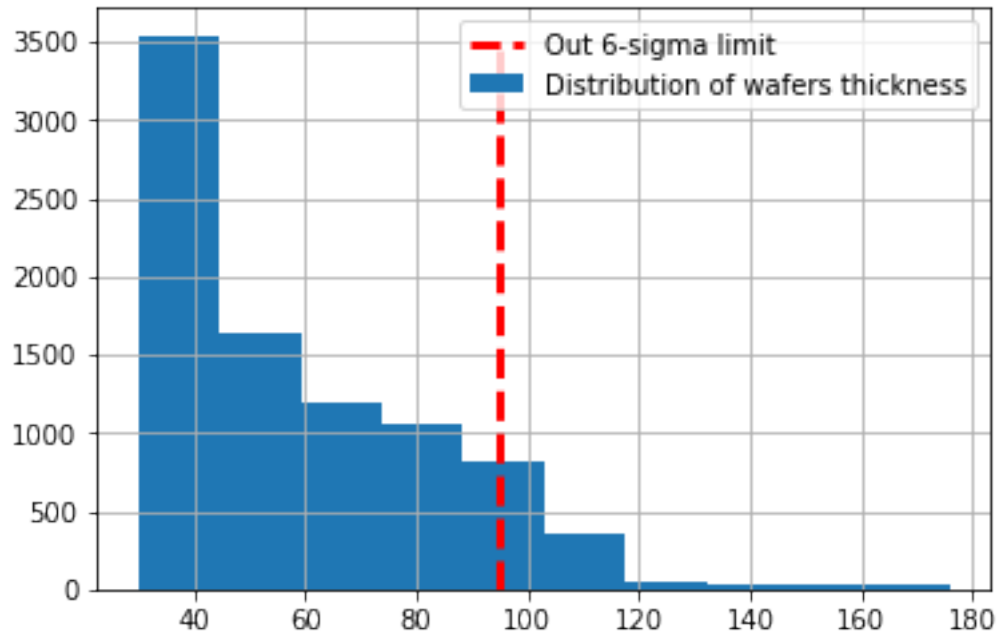
	SawStep	LastEmptying	OutSixSigma	WafersProduction	SiliciumType
0	0	01/01/2017	0	1000	Mono
1	0	01/01/2017	0	3000	Mono
2	0	01/01/2017	0	5000	Mono
3	0	01/01/2017	0	6000	Mono
4	0	01/01/2017	0	7000	Mono

Let's take a look at the data :

```
In [12]: import matplotlib.pyplot as plt
import seaborn as sns

data_maintenance.TTV.hist(label="Distribution of wafers thickness")
plt.plot([95, 95], [0, 3500],
         color='red',
         linestyle='--',
         linewidth=3,
         label="Out 6-sigma limit")

plt.legend(loc='upper right');
plt.show()
```



As we can see, mostly 15% of the production is out 6-sigma limit. It's a lot and it's a quite waste of ressources since the wafers which doesn't respond to quality criteria are thrown away.

That's why this project of predictive maintenance has for first goal to prevent this drift of non quality production, because the cost of these anomalies are pretty heavy for the firm.

### 1.3 3 - Study of hypothesis

At first, before beginning the step of modeling, we have to check the following hypothesis :

- <li>A - Hypothesis of normality of data</li>
- <li>B - Hypothesis of stationnary time serie</li>
- <li>C - Hypothesis of event without memory (Markov process)</li>

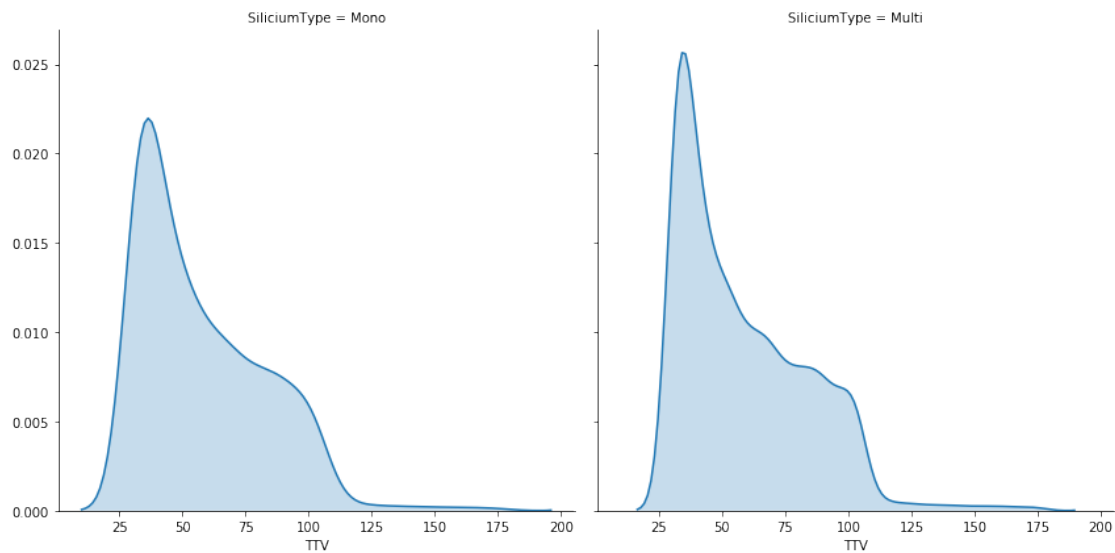
#### 1.3.1 A - Hypothesis of normality of data

The purpose of the verification of this hypothesis is to know if parametric methods can be used for our learning model.

```
In [14]: g = sns.FacetGrid(data_maintenance,
                             col="SiliciumType",
                             size=6)

g.map(sns.kdeplot,
      "TTV",
      shade=True)
```

```
plt.xlabel('TTV')
plt.show()
```



A first visual look would suggest that the parameter observed doesn't follow a gaussian distribution (low Skewness and non symmetric Kurtosis).


We could check also with a QQ plot test (using percentile cumul) :

```
In [15]: import scipy
          from statsmodels.graphics.gofplots import qqplot
          from scipy.stats import shapiro
          from scipy.stats import anderson
          from statsmodels.tsa.stattools import adfuller
          from statsmodels.tsa.stattools import kpss

          from statsmodels.tsa.arima_model import ARIMA

          qqplot(data_maintenance.TTV, line='s')
          plt.show()
```

C:\Users\monne\Anaconda3\lib\site-packages\statsmodels\compat\pandas.py:56: FutureWarning: The from pandas.core import datetools



It seems that the distribution is not a gaussian one...

To be sure, we could compute a statistical test and we will use the following ones : Shapiro Test  
D'Agostino Test  
Anderson Test

```
In [17]: stats,p_value = shapiro(data_maintenance.TTV)
         print('Stats = {} - p-value = {}'.format(stats,p_value))
```

```
Stats = 0.8932368159294128 - p-value = 0.0
```

```
C:\Users\monne\Anaconda3\lib\site-packages\scipy\stats\morestats.py:1310: UserWarning: p-value
  warnings.warn("p-value may not be accurate for N > 5000.")
```

```
In [18]: # D'Agostino and Pearson's Test
         print(scipy.stats.normaltest(data_maintenance.TTV))
```

```
NormaltestResult(statistic=1268.5132451261297, pvalue=3.514379931470798e-276)
```

```
In [28]: # Anderson's Test
         result = anderson(data_maintenance.TTV)
         print('Statistic: %.3f' % result.statistic)
```

```

for i in range(len(result.critical_values)):
    sl, cv = result.significance_level[i], result.critical_values[i]

    if result.statistic < result.critical_values[i]:
        print('%.3f: %.3f, data looks normal (fail to reject H0)' % (sl, cv))
    else:
        print('Critical value of {}% : {} --> data doesn\'t look normal (reject H0)'

Statistic: 259.087
Critical value of 15.0% : 0.576 --> data doesn't look normal (reject H0)
Critical value of 10.0% : 0.656 --> data doesn't look normal (reject H0)
Critical value of 5.0% : 0.787 --> data doesn't look normal (reject H0)
Critical value of 2.0% : 0.918 --> data doesn't look normal (reject H0)
Critical value of 1.0% : 1.092 --> data doesn't look normal (reject H0)

```

## 1.4 B - Hypothesis of stationnary time serie

A first look on the distribution of one of the parameter (the BOW measure which is globally the same as the TTV criteria) would indicate that fluctuations are following a time serie, as we could read on the below picture :

Let's check this assumption :

```

<li>Transform the data</li>
<li>Test of Ad Fuller</li>
<li>ARIMA performing</li>

```

### 1.4.1 Transform the data :

At first, we have to group the data on the variable Date in order to be able to compute the analysis on a daily basis.

```

In [30]: data_maintenance['year'] = data_maintenance['Date'].apply(lambda x:x[len(x)-4:])
        data_maintenance['month'] = data_maintenance['Date'].apply(lambda x:x[len(x)-7:len(x)-4])
        data_maintenance['day'] = data_maintenance['Date'].apply(lambda x:x[:2])

        data_maintenance['date_py'] = pd.to_datetime(data_maintenance[['year', 'month', 'day']])

        grouped = data_maintenance.groupby('date_py')
        mean_TTV = grouped['TTV'].agg(np.mean)
        print(mean_TTV.head(5))

```

```


date_py
2017-01-01    52.333333
2017-01-02    56.291667
2017-01-03    57.291667
2017-01-04    57.333333
2017-01-05    57.333333

```

Name: TTV, dtype: float64

The dataset is now regrouped by the date of the day (we keep the mean of each variable)

```
In [32]: mean_TTV.plot()  
plt.show()
```



The graphical visualisation could let think that we are facing a stationnary time series use case...

Stationnary, because there is no trend, the variance doesn't grow with the date and there is not sign of covariance with the time axis...

#### 1.4.2 Augmented Dickey Fuller Test

```
In [33]: print('--'*20)  
print('Augmented Dickey Fuller Test')  
  
results = adfuller(mean_TTV,  
                    autolag='AIC')  
  
test_statistics,p_value = results[0],results[1]
```

```

print('valeur du test : {} - Valeur de la p-value : {}'.format(test_statistics,
                                                                p_value))

for key,value in results[4].items():
    print('Critical value à {} % : {}'.format(key,value))
print('--'*20)
if p_value < 0.05:
    print('Reject H0 : the serie of fluctuations of TTV is stationnary...')
else:
    print('Fail to reject H0 : the serie of fluctuations of TTV is not stationnary...')

-----
Augmented Dickey Fuller Test
valeur du test : -18.501751389268442 - Valeur de la p-value : 2.1201245071674424e-30
Critical value à 1% % : -3.4483935212959844
Critical value à 5% % : -2.8694912343676497
Critical value à 10% % : -2.571005879151811
-----
Reject H0 : the serie of fluctuations of TTV is stationnary...

```

### 1.4.3 Kwiatkowski-Phillips-Schmidt-Shin Test

```

In [34]: print('--'*20)
         print('Kwiatkowski-Phillips-Schmidt-Shin Test')

results = kpss(mean_TTV,
               regression='c')

test_statistics,p_value = results[0],results[1]

print('valeur du test : {} - Valeur de la p-value : {}'.format(test_statistics,
                                                                p_value))

for key,value in results[3].items():
    print('Critical value à {} % : {}'.format(key,value))
print('--'*20)
if p_value > 0.05:
    print('Fail to reject H0 : the serie of fluctuations of TTV is stationnary...')
else:
    print('Reject H0 : the serie of fluctuations of TTV is not stationnary...')

-----
Kwiatkowski-Phillips-Schmidt-Shin Test
valeur du test : 0.036624491810607694 - Valeur de la p-value : 0.1

```



Critical value à 10% % : 0.347  
Critical value à 5% % : 0.463  
Critical value à 2.5% % : 0.574  
Critical value à 1% % : 0.739

-----  
Fail to reject H0 : the serie of fluctuations of TTV is stationnary...


C:\Users\monne\Anaconda3\lib\site-packages\statsmodels\tsa\stattools.py:1260: InterpolationWarning: warn("p-value is greater than the indicated p-value", InterpolationWarning)

#### 1.4.4 ARIMA performing

```
In [35]: # Forecast with ARIMA model
model = ARIMA(mean_TTV, order=(2,0, 0))
results_ARIMA = model.fit()

plt.plot(mean_TTV)
plt.plot(results_ARIMA.fittedvalues, color='red')
plt.title('RSS: %.4f'% sum((results_ARIMA.fittedvalues-mean_TTV)**2))

Out[35]: Text(0.5,1,'RSS: 53269.1851')
```



The ARIMA model seems specifically poor on this case.

- <li> The predictions (in red) is not aligned on the time series.</li>
- <li> And the RSS score is pretty high... but why ???</li>

After a lot of research and discussion with the process engineers, it seems that data are not correlated with date or time, but with the number of wafers produced.

If nothing is produced while a month, all parameters will remain stable and no drift will be observed. So, our data are not a stationnary time serie use case...

It's a bit tricky since results of tests and visual graphics insinuate the opposite, but this is an effect of an hidden correlation between the date and the cumul of number of wafers produced during the interval...

#### 1.4.5 C - Hypothesis of event without memory (Markov process)

If events observed were without memory, no links would be notified with a past parameter. But on the below visualisation, we could see that fluctutations seem correlated with the past cumul of wafers :

Futhermore, we have observed that drifts are sometimes concomitant with operations of emptying on the saw equipment (oil lubricant). It seems that probability of drift occurence is proportional with the past delay of the last action of emptying performed on the saw equipment :  
#### So, we can deduct that we are not facing a Markov process use case...

## 2 4 - Modeling and learning

- <li>A - Design the features</li>
- <li>B - Select the features</li>
- <li>C - Build the model</li>

### 2.0.1 A - Design the features

in connection with the observations made previously, we will create 4 new features :

- <li>A feature to indicate if it is a day when an emptying action is performed</li>
- <li>A feature to calculate the number of wafers produced since the last day of emptying</li>
- <li>A feaure to count the number of days since the last action of emptying</li>
- <li>A dummy feature for coding the type of silicium material (has an impact on the saw wire at

```
In [43]: is_emptying = list()
        cumul_wafers = list()
        day_since_emptying = list()
        cumul_ant=0

        for index,row in data_maintenance.iterrows():
            if row['Date'][:2]=='01':
                is_emptying.append(1)
                cumul_wafers.append(0)
                cumul_ant=row['WafersProduction']
                day=int(row['LastEmptying'][:2])
                day_since_emptying.append(0)
```

```

else:
    is_emptying.append(0)
    cumul_wafers.append(row['WafersProduction']-cumul_ant)
    day_since_emptying.append(np.abs(int(row['Date'][:2])-day))

data_maintenance['day of emptying'] = is_emptying
data_maintenance['cumul of wafers'] = cumul_wafers
data_maintenance['day since emptying'] = day_since_emptying
data_maintenance['Is Mono'] = data_maintenance['SiliciumType'].apply(lambda x: 1 if x

```

```

In [47]: data_maintenance[['Date','day of emptying','cumul of wafers','LastEmptying','day since

```

```

Out[47]:
      Date  day of emptying  cumul of wafers  LastEmptying \
8761  01/01/2018           1              0    01/01/2018
8762  01/01/2018           1              0    01/01/2018
8763  01/01/2018           1              0    01/01/2018

      day since emptying  Is Mono
8761                   0        0
8762                   0        0
8763                   0        0

```

## 2.0.2 B - Select the features


At first, we are checking the correlation between variables since models are sometimes sensitive to it :

```

In [48]: plt.matshow(data_maintenance.corr())

Out[48]: <matplotlib.image.AxesImage at 0x19eb99787b8>

```



```
In [52]: for col in data_maintenance.corr().columns:
          conditions = (data_maintenance.corr()[col]>0.6) & (data_maintenance.corr()[col]<1)
          if len(data_maintenance.corr()[conditions][col])>0:
              columns_correlated = data_maintenance.corr()[conditions][col]
              print('--'*20)
              print(col)
              print('--'*20)
              print(columns_correlated)
```

```
-----
TTV
-----
```

```
OutSixSigma    0.682667
Name: TTV, dtype: float64
-----
```

```
Thickness
-----
```

```
SawStep        0.604948
Name: Thickness, dtype: float64
-----
```

```
SawStep
-----
```

```
Thickness      0.604948
Name: SawStep, dtype: float64
-----
```

```

OutSixSigma
-----
TTV      0.682667
Name: OutSixSigma, dtype: float64
-----
cumul of wafers
-----
day since emptying      0.999174
Name: cumul of wafers, dtype: float64
-----
day since emptying
-----
cumul of wafers      0.999174
Name: day since emptying, dtype: float64

```

**Then we are trying to detect the variables which explain the most of variability in the dataset with a PCA :**

```

In [53]: from sklearn.preprocessing import StandardScaler

features = data_maintenance.corr().columns.tolist()
features.remove('OutSixSigma')
target = 'OutSixSigma'

X = data_maintenance.loc[:,features]
y = data_maintenance.loc[:,target]

X = StandardScaler().fit_transform(X)

from sklearn.decomposition import PCA

pca = PCA(n_components=4)
principal_components = pca.fit_transform(X)

principalDf = pd.DataFrame(data = principal_components,
                           columns = ['PCI' + str(i) for i in range(4)])


plt.scatter(principalDf.iloc[:,0],
            principalDf.iloc[:,1],
            c=y,
            label=y.unique())

print('Variances explained by axes : {}'.format(pca.explained_variance_ratio_))

plt.show()

Variances explained by axes : [ 0.31732599  0.21365304  0.13912937  0.12088666]

```



```
In [54]: sortedFeatures_axe1 = [i[0] for i in sorted(zip(features,pca.components_[0]),
                                                    key=lambda l:l[1],reverse=True)]

sortedFeatures_axe2 = [i[0] for i in sorted(zip(features,pca.components_[1]),
                                                    key=lambda l:l[1],reverse=True)]

sortedFeatures_axe3 = [i[0] for i in sorted(zip(features,pca.components_[2]),
                                                    key=lambda l:l[1],reverse=True)]

sortedFeatures_axe4 = [i[0] for i in sorted(zip(features,pca.components_[3]),
                                                    key=lambda l:l[1],reverse=True)]

print(sortedFeatures_axe1[:1])
print(sortedFeatures_axe2[:1])
print(sortedFeatures_axe3[:1])
print(sortedFeatures_axe4[:1])

['Is Mono']
['day of emptying']
['SawGroove']
['TTV']
```

Then we have our features which will carry the most of variability for our future learning model.

As we can see on the plan projection above, our data are segmented by the type of silicium and the proximity of the day of emptying.

It seems that the day of emptying has an impact on the target feature ("OutOf6Sigma") since points in yellow on the plan, are systemtically above the others (when the axis vertical is carried by the "day of emptying" feature).

### 2.0.3 C - Build the model

```
In [55]: from sklearn.model_selection import train_test_split

train_features, test_features, train_target, test_target = train_test_split(X,
                                                                              y,
                                                                              test_size=0.2,
                                                                              stratify=y)

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import auc, roc_curve

model_regression = LogisticRegression()
model_regression.fit(train_features, train_target)
predictions = model_regression.predict(test_features)
probalities = model_regression.predict_proba(test_features)

accuracy = predictions == test_target

true_positives = (predictions==1) & (test_target==1)
false_positives = (predictions==1) & (test_target==0)
true_negatives = (predictions==0) & (test_target==0)
false_negatives = (predictions==0) & (test_target==1)

sensitivity = np.sum(true_positives) / (np.sum(true_positives)+np.sum(false_negatives))

specificity = np.sum(true_negatives) / (np.sum(true_negatives) + np.sum(false_positives))


print('Sensibility of predictions is about : {}'.format(sensitivity))
print('Specificity of predictions is about : {}'.format(specificity))

fpr, tpr, thresholds = roc_curve(test_target, probalities[:,1], pos_label=1)
auc = auc(fpr, tpr)

print('Perimeter of AUC curb : {}'.format(auc))

plt.plot(fpr, tpr, label='AUC={}'.format(auc))
plt.plot([0,1],[0,1], color='red')
plt.legend(loc='lower right')
plt.show()
```

Sensibility of predictions is about : 1.0  
Specificity of predictions is about : 1.0  
Perimeter of AUC curb : 1.0



The score of accuracy seems perfect... too perfect. When performing this logistic regression, we made a rookie mistake. We keep features which contain information on the future... the measure of thickness or TTV are not known when the cut is performed... these values are measured a few hours later by a specific equipment. It a bias...

When the prediction will be performed, no information about thickness of the wafer will be available. That's why we have to remove these features to build our model :


```
In [59]: X_reduced = ['cumul of wafers',  
                      'day since emptying',  
                      'Is Mono',  
                      'day of emptying',  
                      'WafersProduction']  
  
X = data_maintenance.loc[:,X_reduced]  
y = data_maintenance.loc[:,target]  
  
X = StandardScaler().fit_transform(X)  
  
train_features, test_features, train_target, test_target = train_test_split(X,
```



```
y,  
test_size=  
stratify=
```

```
from sklearn.metrics import auc,roc_curve  
  
penalty = {  
    0:10,  
    1:1  
}  
  
model_regression = LogisticRegression(class_weight=penalty)  
  
model_regression.fit(train_features,train_target)  
predictions = model_regression.predict(test_features)  
probalities = model_regression.predict_proba(test_features)  
  
accuracy = predictions == test_target  
  
true_positives = (predictions==1) & (test_target==1)  
false_positives = (predictions==1) & (test_target==0)  
true_negatives = (predictions==0) & (test_target==0)  
false_negatives = (predictions==0) & (test_target==1)  
  
sensitivity = np.sum(true_positives) / (np.sum(true_positives)+np.sum(false_negatives))  
  
specificity = np.sum(true_negatives) / (np.sum(true_negatives) + np.sum(false_positives))  
  
print('Sensibility of predictions is about : {}'.format(sensitivity))  
print('Specificity of predictions is about : {}'.format(specificity))  
  
fprs, tprs, thresholds = roc_curve(test_target, probalities[:,1], pos_label=1)  
aucs = auc(fprs,tprs)  
  
print('Perimeter of AUC curb : {}'.format(aucs))  
  
plt.plot(fprs,tprs,label='AUC={}'.format(aucs))  
plt.plot([0,1],[0,1],color='red')  
plt.legend(loc='lower right')  
plt.show()
```

```
Sensibility of predictions is about : 0.20491803278688525  
Specificity of predictions is about : 1.0  
Perimeter of AUC curb : 0.6286194818425994
```



Right now, the accuracy of the model is less impressive. The cuts with drifts 'out of 6 sigma specifications' are very poorly predicted (about 20%) The only good point is that the model never made a prediction of non drift which was false. But, this is not sufficient and we have to try better model than logistic regression...

We will now perform an ensembled model called "Random Forest". Ensembled since many instances of decision trees are running to compute a accurate prediction on the binary target.

```
In [64]: from sklearn.ensemble import RandomForestClassifier
         from sklearn.metrics import auc,roc_curve

         rf = RandomForestClassifier(n_estimators=10,
                                   random_state=1,
                                   class_weight='balanced')

         rf.fit(train_features,train_target)
         predictions = rf.predict(test_features)
         probabilities = rf.predict_proba(test_features)

         accuracy = predictions == test_target

         true_positives = (predictions==1) & (test_target==1)
         false_positives = (predictions==1) & (test_target==0)
         true_negatives = (predictions==0) & (test_target==0)
         false_negatives = (predictions==0) & (test_target==1)
```

```

sensitivity = np.sum(true_positives) / (np.sum(true_positives)+np.sum(false_negatives))

specificity = np.sum(true_negatives) / (np.sum(true_negatives) + np.sum(false_positives))

print('Sensibility of predictions is about : {}'.format(sensitivity))
print('Specificity of predictions is about : {}'.format(specificity))


fprs, tprs, thresholds = roc_curve(test_target, probabilities[:,1], pos_label=1)
aucs = auc(fprs,tprs)

print('Perimeter of AUC curb : {}'.format(aucs))

plt.plot(fprs,tprs,label='AUC={}'.format(aucs))
plt.plot([0,1],[0,1],color='red')
plt.legend(loc='lower right')
plt.show()

```

Sensibility of predictions is about : 0.27049180327868855  
 Specificity of predictions is about : 0.9774011299435028  
 Perimeter of AUC curb : 0.768103724098445



The results are better so we can attempt a first move with this model...

## 2.1 5 - Results and operation

In production, an input of values of the features are sent to the model before the ingot been cut in the saw equipment. If the model predict a drift, an operation of preventive maintenance (change of diamant cutwire for example is then performed in order to prevent the drift or the anomaly. And the ingot is cut into wafers...

```
In [67]: def prediction_maintenance(cumul_wafers,days_since_emptying,is_mono,days_of_emptying,WaferProduction)

    data = [cumul_wafers,
            days_since_emptying,
            is_mono,
            days_of_emptying,
            WaferProduction]

    data_res = np.asarray(data).reshape(1,-1)

    #data = StandardScaler().fit_transform(data_res)
    data = data_res
    prevision = rf.predict(data)
    prevision_proba = rf.predict_proba(data)

    if prevision:
        print('With the data given in input, we can predict that a maintenance action is needed')
    else :
        print('With the data given in input, no need at the moment for a maintenance action')

    print('Probabilities of need of action maintenance {}'.format(prevision_proba[:,1]))

    return prevision_proba

#
print('--'*20)
### day 1
cumul_wafers = 10
days_since_emptying = 1
is_mono = 0
days_of_emptying= 0
WaferProduction= 0

prediction = prediction_maintenance(cumul_wafers,days_since_emptying,is_mono,days_of_emptying,WaferProduction)
print('--'*20)

### day 24
cumul_wafers = 879000
days_since_emptying = 24
is_mono = 1
days_of_emptying= 0
WaferProduction= 8790000
```

```

prediction = prediction_maintenance(cumul_wafers,days_since_emptying,is_mono,days_of_c
print('--'*20)

```

-----  
 With the data given in input, we can predict that a maintenance action is necessary !!!  
 Probabilities of need of action maintenance [ 70.]%  
 -----

With the data given in input, no need at the moment for a maintenance action...  
 Probabilities of need of action maintenance [ 30.]%  
 -----


**The model is perfectible but it gives some interesting results. For example, the ratio of non quality on wafers cut dropped at 5% when it culminates at 15% before the predictive model.**

```

In [69]: data_maintenance.TTV.hist(label="Distribution of wafers thickness",color="green")
plt.plot([115, 115], [0, 3500],
         color='red',
         linestyle='--',
         linewidth=3,
         label="Out 6-sigma limit")

plt.legend(loc='upper right');
plt.show()

```



**Then, this project was a good example of a smart use of data science for accurate prediction of action maintenance... A first step towards industry 4.0 for solar manufacturing !!!**