

Spark SQL

August 22, 2018

1 Spark SQL

```
In [ ]: # Find path to PySpark.  
import findspark  
findspark.init()
```

```
In [2]: # Import PySpark and initialize SparkContext object.  
import pyspark  
sc = pyspark.SparkContext()
```

```
In [10]: # Read `recent-grads.csv` in to an RDD.  
path = 'E:\\OneDrive\\Formation\\2018 - DataQuest\\Codes\\_Solutions-master'  
file = 'recent-grads.csv'  
file_path = '\\'.join([path,file])  
  
f = sc.textFile(file_path)
```

1.1 RDD Transformations

```
In [11]: # Slice data in lines  
data = f.map(lambda line: line.split('\n'))
```

```
In [6]: data.take(5)
```

```
Out[6]: [['Rank,Major_code,Major>Total,Men,Women,Major_category,ShareWomen,Sample_size,Employed  
[1,2419,PETROLEUM ENGINEERING,2339,2057,282,Engineering,0.120564344,36,1976,1849,270  
[2,2416,MINING AND MINERAL ENGINEERING,756,679,77,Engineering,0.101851852,7,640,556,  
[3,2415,METALLURGICAL ENGINEERING,856,725,131,Engineering,0.153037383,3,648,558,133,  
[4,2417,NAVAL ARCHITECTURE AND MARINE ENGINEERING,1258,1123,135,Engineering,0.107313,
```

```
In [21]: s = ['2,2416,MINING AND MINERAL ENGINEERING,756,679,77,Engineering,0.101851852,7,640,  
for i in s:  
    print(i)
```

```
2,2416,MINING AND MINERAL ENGINEERING,756,679,77,Engineering,0.101851852,7,640,556,170,388,85,
```

```
In [22]: s[0]
```

```
Out[22]: '2,2416,MINING AND MINERAL ENGINEERING,756,679,77,Engineering,0.101851852,7,640,556,1'
```

```
In [25]: def Search_Job(line):
          id = line[0]
          if 'ENGINEERING' in line[0]:
              yield id, 'Good Job : '

          good_jobs = data.flatMap(lambda x:Search_Job(x))
```

```
In [26]: good_jobs.take(5)
```

```
Out[26]: [('1,2419,PETROLEUM ENGINEERING,2339,2057,282,Engineering,0.120564344,36,1976,1849,270,
          'Good Job : '),
          ('2,2416,MINING AND MINERAL ENGINEERING,756,679,77,Engineering,0.101851852,7,640,556,1
          'Good Job : '),
          ('3,2415,METALLURGICAL ENGINEERING,856,725,131,Engineering,0.153037383,3,648,558,133
          'Good Job : '),
          ('4,2417,NAVAL ARCHITECTURE AND MARINE ENGINEERING,1258,1123,135,Engineering,0.10731
          'Good Job : '),
          ('5,2405,CHEMICAL ENGINEERING,32260,21239,11021,Engineering,0.341630502,289,25694,23
          'Good Job : ')]
```

```
In [27]: def naval_job(line):
          if 'naval' in line[0].lower():
              return True
          else:
              return False

          sea_jobs = data.filter(lambda x:naval_job(x))
```

```
In [28]: sea_jobs.take(5)
```

```
Out[28]: [['4,2417,NAVAL ARCHITECTURE AND MARINE ENGINEERING,1258,1123,135,Engineering,0.10731
```

1.2 RDD Actions

```
In [29]: data_count = data.count()
          data_count
```

```
Out[29]: 174
```

```
In [31]: data_collect = data.collect()[:3]
          data_collect
```

```
Out[31]: [['Rank,Major_code,Major>Total,Men,Women,Major_category,ShareWomen,Sample_size,Employ
          ['1,2419,PETROLEUM ENGINEERING,2339,2057,282,Engineering,0.120564344,36,1976,1849,270
          ['2,2416,MINING AND MINERAL ENGINEERING,756,679,77,Engineering,0.101851852,7,640,556
```

1.3 Transforming Dataset

```
In [43]: def return_id(line):
          data = line[0]
          if data[0]!='R':
              raw_data = data.split(',')
              return int(raw_data[0])

          data_ids = data.map(lambda line:return_id(line))

In [45]: list_data_ids = data_ids.collect()
          list_data_ids[:5]

Out[45]: [None, 1, 2, 3, 4]

In [46]: real_text = data_ids.filter(lambda x:x is not None)

          list_data_ids = real_text.collect()
          list_data_ids[:5]

Out[46]: [1, 2, 3, 4, 5]
```

1.4 Spark DataFrames

```
In [62]: import csv
          with open(file_path,'r') as f:
              body=list(csv.reader(f))

In [65]: header = body[:1][0]
          body = body[1:]

In [47]: # Import SQLContext
          from pyspark.sql import SQLContext

          # Pass in the SparkContext object `sc`
          sqlCtx = SQLContext(sc)

In [68]: df = sqlCtx.read.csv(file_path,
                               header=True)

In [69]: print(type(df))

<class 'pyspark.sql.dataframe.DataFrame'>

In [50]: df.printSchema()
```

```

root
|-- _c0: string (nullable = true)
|-- _c1: string (nullable = true)
|-- _c2: string (nullable = true)
|-- _c3: string (nullable = true)
|-- _c4: string (nullable = true)
|-- _c5: string (nullable = true)
|-- _c6: string (nullable = true)
|-- _c7: string (nullable = true)
|-- _c8: string (nullable = true)
|-- _c9: string (nullable = true)
|-- _c10: string (nullable = true)
|-- _c11: string (nullable = true)
|-- _c12: string (nullable = true)
|-- _c13: string (nullable = true)
|-- _c14: string (nullable = true)
|-- _c15: string (nullable = true)
|-- _c16: string (nullable = true)
|-- _c17: string (nullable = true)
|-- _c18: string (nullable = true)
|-- _c19: string (nullable = true)
|-- _c20: string (nullable = true)

```

```
In [70]: df.show(5)
```

Rank	Major_code	Major	Total	Men	Women	Major_category	ShareWomen	Sample_size
1	2419	PETROLEUM ENGINEE...	2339	2057	282	Engineering	0.120564344	36
2	2416	MINING AND MINERA...	756	679	77	Engineering	0.101851852	7
3	2415	METALLURGICAL ENG...	856	725	131	Engineering	0.153037383	3
4	2417	NAVAL ARCHITECTUR...	1258	1123	135	Engineering	0.107313196	16
5	2405	CHEMICAL ENGINEERING	32260	21239	11021	Engineering	0.341630502	289

only showing top 5 rows

```
In [71]: df.select('Major', 'Total', 'Men', 'Women').show()
```

Major	Total	Men	Women
PETROLEUM ENGINEE...	2339	2057	282
MINING AND MINERA...	756	679	77
METALLURGICAL ENG...	856	725	131
NAVAL ARCHITECTUR...	1258	1123	135

CHEMICAL ENGINEERING	32260	21239	11021
NUCLEAR ENGINEERING	2573	2200	373
ACTUARIAL SCIENCE	3777	2110	1667
ASTRONOMY AND AST...	1792	832	960
MECHANICAL ENGINE...	91227	80320	10907
ELECTRICAL ENGINE...	81527	65511	16016
COMPUTER ENGINEERING	41542	33258	8284
AEROSPACE ENGINEE...	15058	12953	2105
BIOMEDICAL ENGINE...	14955	8407	6548
MATERIALS SCIENCE	4279	2949	1330
ENGINEERING MECHA...	4321	3526	795
BIOLOGICAL ENGINE...	8925	6062	2863
INDUSTRIAL AND MA...	18968	12453	6515
GENERAL ENGINEERING	61152	45683	15469
ARCHITECTURAL ENG...	2825	1835	990
COURT REPORTING	1148	877	271

+-----+-----+-----+-----+

only showing top 20 rows

```
In [77]: Major_2000 = df[df['Total']>2000]
        Major_2000.select('Major', 'Total', 'Men', 'Women').show(5)
```

+-----+-----+-----+-----+			
Major Total Men Women			
+-----+-----+-----+-----+			
PETROLEUM ENGINEE...	2339	2057	282
CHEMICAL ENGINEERING	32260	21239	11021
NUCLEAR ENGINEERING	2573	2200	373
ACTUARIAL SCIENCE	3777	2110	1667
MECHANICAL ENGINE...	91227	80320	10907

+-----+-----+-----+-----+

only showing top 5 rows

```
In [83]: Female = df[df['Women']>=df['Men']].select('Major', 'Total', 'Men', 'Women').show(5)
        Female
```

+-----+-----+-----+-----+			
Major Total Men Women			
+-----+-----+-----+-----+			
PETROLEUM ENGINEE...	2339	2057	282
MINING AND MINERA...	756	679	77
NAVAL ARCHITECTUR...	1258	1123	135
NUCLEAR ENGINEERING	2573	2200	373
ASTRONOMY AND AST...	1792	832	960

+-----+-----+-----+-----+

only showing top 5 rows

```
In [86]: pandas_df = df.toPandas()
```

```
In [99]: def Men_Women(line):
        if line['Women'] and line['Men']:
            if line['Women']>=line['Men']:
                return True
            else:
                return False
        else:
            return False
```

```
pandas_df = pandas_df.dropna()
pandas_df['Egalitiy']=pandas_df.apply(Men_Women,axis=1)
```

```
In [100]: pandas_df_women = pandas_df.where(pandas_df['Egalitiy']==True)
pandas_df_women.head(3)
```

```
Out[100]:
```

	Rank	Major_code		Major	Total	Men	Women	\
0	1	2419		PETROLEUM ENGINEERING	2339	2057	282	
1	2	2416		MINING AND MINERAL ENGINEERING	756	679	77	
2	NaN	NaN		NaN	NaN	NaN	NaN	

	Major_category	ShareWomen	Sample_size	Employed	...	\
0	Engineering	0.120564344	36	1976	...	
1	Engineering	0.101851852	7	640	...	
2	NaN	NaN	NaN	NaN	...	


	Full_time_year_round	Unemployed	Unemployment_rate	Median	P25th	P75th	\
0	1207	37	0.018380527	110000	95000	125000	
1	388	85	0.117241379	75000	55000	90000	
2	NaN	NaN	NaN	NaN	NaN	NaN	

	College_jobs	Non_college_jobs	Low_wage_jobs	Egalitiy
0	1534	364	193	1.0
1	350	257	50	1.0
2	NaN	NaN	NaN	NaN

[3 rows x 22 columns]

```
In [101]: from matplotlib import pyplot as plt
          %matplotlib inline
```

```
pandas_df['Total'].hist()
plt.show()
```



1.5 Spark SQL

```
In [110]: # Import SQLContext
          from pyspark.sql import SQLContext

          # Pass in the SparkContext object `sc`
          sqlCtx = SQLContext(sc)

In [111]: df = sqlCtx.read.csv(file_path,
                                header=True)

In [112]: df.registerTempTable('major_results')

          tables = sqlCtx.tableNames()

          print(tables)

['major_results']

In [105]: query = "SELECT women FROM major_results"
          sqlCtx.sql(query).show()

+-----+
|women|
+-----+
```

```
| 282|
| 77|
| 131|
| 135|
|11021|
| 373|
| 1667|
| 960|
|10907|
|16016|
| 8284|
| 2105|
| 6548|
| 1330|
| 795|
| 2863|
| 6515|
|15469|
| 990|
| 271|
+-----+
```

only showing top 20 rows

```
In [116]: query = """SELECT men,women FROM major_results"""
```

```
df = sqlCtx.sql(query)
```

```
df.describe().show()
```

```
Out[116]: DataFrame[women: string]
```

```
In [ ]: df1 = sqlCtx.read.csv(file_path,
                             header=True)
```

```
df2 = sqlCtx.read.csv(file_path,
                     header=True)
```

```
tables = sqlCtx.tableNames()
```

```
In [ ]: df1.registerTempTable('major1')
df2.registerTempTable('major2')
```

```
query = """SELECT
            t1.total,
            t2.total
        FROM
            major1 as t1
```



```

        INNER JOIN
            major2 as t2
        ON t1.major_code = t2.major_code
    """

In [ ]: sqlCtx.sql(query).show()

In [ ]: df1.registerTempTable('major1')
        df2.registerTempTable('major2')

        query = """SELECT
            SUM(t1.total),
            SUM(t2.total)
        FROM
            major1 as t1
        INNER JOIN
            major2 as t2
        ON t1.major_code = t2.major_code
    """

In [ ]: sqlCtx.sql(query).show()

```